

# Week 5: CNN-1

## 实验准备

- 熟悉python语言的使用和numpy,torch的基本用法
- 熟悉神经网络的训练过程与优化方法
- 结合理论课的内容,了解卷积与卷积神经网络(CNN)的内容和原理
- 了解常用的CNN模型的基本结构,如AlexNet,Vgg,ResNet

## 实验过程


### 1. 卷积与卷积层

- numpy实现卷积
- pytorch中的卷积层和池化层

### 2. CNN

- 实现并训练一个基本的CNN网络
- ResNet
- VGG

卷积



在实验课上我们已经了解过卷积运算的操作当我们对一张二维的图像做卷积时,将卷积核沿着图像进行滑动乘加即可(如上图所示).

下面的conv函数实现了对二维单通道图像的卷积.考虑输入的卷积核kernel的长宽相同,padding为对图像的四个边缘补0 stride为卷积核窗口滑动的步长

In [1]:

```
import numpy as np

def convolution(img, kernel, padding=1, stride=1):
    """
    img: input image with one channel
    kernel: convolution kernel
    """

    h, w = img.shape
    kernel_size = kernel.shape[0]

    # height and width of image with padding
    ph, pw = h + 2 * padding, w + 2 * padding
    padding_img = np.zeros((ph, pw))
    padding_img[padding:h + padding, padding:w + padding] = img

    # height and width of output image
    result_h = (h + 2 * padding - kernel_size) // stride + 1
    result_w = (w + 2 * padding - kernel_size) // stride + 1

    result = np.zeros((result_h, result_w))

    # convolution
    x, y = 0, 0
    for i in range(0, ph - kernel_size + 1, stride):
        for j in range(0, pw - kernel_size + 1, stride):
            roi = padding_img[i:i+kernel_size, j:j+kernel_size]
            result[x, y] = np.sum(roi * kernel)
            y += 1
        y = 0
        x += 1
    return result
```

下面在图像上简单一下测试我们的conv函数,这里使用3\*3的高斯核对下面的图像进行滤波.



In [2]:

```
from PIL import Image
import matplotlib.pyplot as plt
img = Image.open('pics/lena.jpg').convert('L')
plt.imshow(img, cmap='gray')
```

Out[2]:

```
<matplotlib.image.AxesImage at 0x7f99ba7f75c0>
```

In [3]:

```

# a Laplace kernel
laplace_kernel = np.array([[ -1, -1, -1],
                           [-1, 8, -1],
                           [-1, -1, -1]])

# Gauss kernel with kernel_size=3
gauss_kernel3 = (1/ 16) * np.array([[1, 2, 1],
                                    [2, 4, 2],
                                    [1, 2, 1]])

# Gauss kernel with kernel_size=5
gauss_kernel5 = (1/ 84) * np.array([[1, 2, 3, 2, 1],
                                    [2, 5, 6, 5, 2],
                                    [3, 6, 8, 6, 3],
                                    [2, 5, 6, 5, 2],
                                    [1, 2, 3, 2, 1]])

fig, ax = plt.subplots(1, 3, figsize=(12, 8))

laplace_img = convolution(np.array(img), laplace_kernel, padding=1, stride=1)
ax[0].imshow(Image.fromarray(laplace_img), cmap='gray')
ax[0].set_title('laplace')

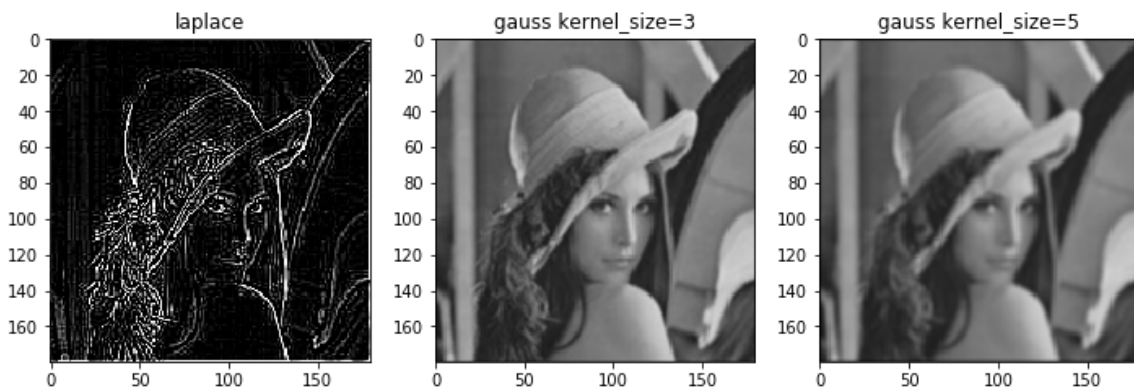
gauss3_img = convolution(np.array(img), gauss_kernel3, padding=1, stride=1)
ax[1].imshow(Image.fromarray(gauss3_img), cmap='gray')
ax[1].set_title('gauss kernel_size=3')

gauss5_img = convolution(np.array(img), gauss_kernel5, padding=2, stride=1)
ax[2].imshow(Image.fromarray(gauss5_img), cmap='gray')
ax[2].set_title('gauss kernel_size=5')

```

Out[3]:

Text(0.5,1,'gauss kernel\_size=5')



上面我们实现了实现了对单通道输入单通道输出的卷积.在CNN中,一般使用到的都是多通道输入多通道输出的卷积,要实现多通道的卷积,我们只需要对循环调用上面的conv函数即可.

In [4]:

```
def myconv2d(features, weights, padding=0, stride=1):
    """
    features: input, in_channel * h * w
    weights: kernel, out_channel * in_channel * kernel_size * kernel_size
    return output with out_channel
    """
    in_channel, h, w = features.shape
    out_channel, _, kernel_size, _ = weights.shape

    # height and width of output image
    output_h = (h + 2 * padding - kernel_size) // stride + 1
    output_w = (w + 2 * padding - kernel_size) // stride + 1
    output = np.zeros((out_channel, output_h, output_w))

    # call convolution out_channel * in_channel times
    for i in range(out_channel):
        weight = weights[i]
        for j in range(in_channel):
            feature_map = features[j]
            kernel = weight[j]
            output[i] += convolution(feature_map, kernel, padding, stride)
    return output
```

接下来, 让我们测试我们写好的myconv2d函数.

In [5]:

```

input_data=[
    [[0,0,2,2,0,1],
     [0,2,2,0,0,2],
     [1,1,0,2,0,0],
     [2,2,1,1,0,0],
     [2,0,1,2,0,1],
     [2,0,2,1,0,1]],

    [[2,0,2,1,1,1],
     [0,1,0,0,2,2],
     [1,0,0,2,1,0],
     [1,1,1,1,1,1],
     [1,0,1,1,1,2],
     [2,1,2,1,0,2]]
]
weights_data=[[
    [ 0, 1, 0],
    [ 1, 1, 1],
    [ 0, 1, 0]],

    [[-1, -1, -1],
     [-1, 8, -1],
     [-1, -1, -1]]
]]

# numpy array
input_data = np.array(input_data)
weights_data = np.array(weights_data)

# show the result
print(myconv2d(input_data, weights_data, padding=3, stride=3))

[[[ 0.  0.  0.  0.]
  [ 0.  8. 10.  0.]
  [ 0. -5.  2.  0.]
  [ 0.  0.  0.  0.]]]

```

在Pytorch中,已经为我们提供了卷积和卷积层的实现.使用同样的input和weights,以及stride,padding,pytorch的卷积的结果应该和我们的一样.可以在下面的代码中进行验证.

In [6]:

```

import torch
import torch.nn.functional as F
input_tensor = torch.tensor(input_data).unsqueeze(0).float()
F.conv2d(input_tensor, weight=torch.tensor(weights_data).float(), bias=None, stride=3, padding=3)

```

Out[6]:

```

tensor([[[[ 0.,  0.,  0.,  0.],
           [ 0.,  8., 10.,  0.],
           [ 0., -5.,  2.,  0.],
           [ 0.,  0.,  0.,  0.]]]])

```

## 作业:

上述代码中convolution的实现只考虑卷积核以及padding和stride长宽一致的情况,若输入的卷积核可能长宽不一致,padding与stride的输入可能为两个元素的元祖(代表两个维度上的padding与stride)并使用下面test input对你的convolutionV2进行测试.

In [7]:

```
def convolutionV2(img, kernel, padding=(0,0), stride=(1,1)):
    """
    img: input image with one channel
    kernel: convolution kernel
    """

    h, w = img.shape
    kernel_size_h, kernel_size_w = kernel.shape
    padding_h, padding_w = padding[0], padding[1]
    stride_h, stride_w = stride[0], stride[1]

    # height and width of image with padding
    ph, pw = h + 2 * padding_h, w + 2 * padding_w
    padding_img = np.zeros((ph, pw))
    padding_img[padding_h:h + padding_h, padding_w:w + padding_w] = img

    # height and width of output image
    result_h = (h + 2 * padding_h - kernel_size_h) // stride_h + 1
    result_w = (w + 2 * padding_w - kernel_size_w) // stride_w + 1

    result = np.zeros((result_h, result_w))

    # convolution
    x, y = 0, 0
    for i in range(0, ph - kernel_size_h + 1, stride_h):
        for j in range(0, pw - kernel_size_w + 1, stride_w):
            roi = padding_img[i:i+kernel_size_h, j:j+kernel_size_w]
            result[x, y] = np.sum(roi * kernel)
            y += 1
        y = 0
        x += 1
    return result
```

In [8]:

```
# test input
test_input = np.array([[1, 1, 2, 1],
                        [0, 1, 0, 2],
                        [2, 2, 0, 2],
                        [2, 2, 2, 1],
                        [2, 3, 2, 3]])

test_kernel = np.array([[1, 0], [0, 1], [0, 0]])

# output
print(convolutionV2(test_input, test_kernel, padding=(1, 0), stride=(1, 1)))
print('\n')
print(convolutionV2(test_input, test_kernel, padding=(2, 1), stride=(1, 2)))
print('\n')
```

```
[[ 1.  2.  1.]
 [ 2.  1.  4.]
 [ 2.  1.  2.]
 [ 4.  4.  1.]
 [ 5.  4.  5.]]
```

```
[[ 0.  0.  0.]
 [ 1.  2.  0.]
 [ 0.  1.  1.]
 [ 2.  1.  2.]
 [ 2.  4.  2.]
 [ 2.  4.  1.]
 [ 0.  3.  3.]]
```

## 卷积层

Pytorch提供了卷积层和池化层供我们使用.

卷积层与上面相似, 而池化层与卷积层相似, Pooling layer的主要目的是缩小features的size.常用的有MaxPool(滑动窗口取最大值)与AvgPool(滑动窗口取均值)

In [9]:

```
import torch
import torch.nn as nn

x = torch.randn(1, 1, 32, 32)

conv_layer = nn.Conv2d(in_channels=1, out_channels=3, kernel_size=3, stride=1, padding=0)
y = conv_layer(x)
print(x.shape)
print(y.shape)

torch.Size([1, 1, 32, 32])
torch.Size([1, 3, 30, 30])
```



请问:

1. 输入与输出的tensor的size分别是多少?该卷积层的参数量是多少?
2. 若kernel\_size=5,stride=2,padding=2, 输出的tensor的size是多少?在上述代码中改变参数后试验后并回答.
3. 若输入的tensor size为N\*C\*H\*W,若第5行中卷积层的参数为  
in\_channels=C,out\_channels=Cout,kernel\_size=k,stride=s,padding=p,那么输出的tensor size是多少?

In [10]:

```
import torch
import torch.nn as nn

x = torch.randn(1, 1, 32, 32)

conv_layer = nn.Conv2d(in_channels=1, out_channels=3, kernel_size=5, stride=2, padding=2)
y = conv_layer(x)
print(x.shape)
print(y.shape)

torch.Size([1, 1, 32, 32])
torch.Size([1, 3, 16, 16])
```

答:

1.  $size_{in} = 32$ ;  $size_{out} = 30$ ;  $F \times F \times C_{input} \times K + K = 3 * 3 * 1 * 3 + 3 = 30$  [Ref. \(https://blog.csdn.net/gaishi\\_hero/article/details/81512404\)](https://blog.csdn.net/gaishi_hero/article/details/81512404)
2.  $size_{out} = 16$ .
3.  $\min((h + 2 * p - k) // s + 1, (w + 2 * p - k) // s + 1)$

In [11]:

```
# input N * C * H * W
x = torch.randn(1, 1, 4, 4)

# maxpool
maxpool = nn.MaxPool2d(kernel_size=2, stride=2)
y = maxpool(x)

# avgpool
avgpool = nn.AvgPool2d(kernel_size=2, stride=2)
z = avgpool(x)

#avgpool
print(x)
print(y)
print(z)

tensor([[[[[-0.7988, -0.6036,  1.0944,  1.0869],
          [ 1.1715, -1.8142, -0.5802,  1.5753],
          [ 1.3232,  0.6413, -0.5604,  0.9052],
          [-0.3123,  1.1715,  0.0411, -0.0606]]]])
tensor([[[[1.1715, 1.5753],
          [1.3232, 0.9052]]]])
tensor([[[[-0.5113,  0.7941],
          [ 0.7059,  0.0813]]]])
```

## GPU

我们可以选择在cpu或gpu上来训练我们的模型.

实验室提供了4卡的gpu服务器,要查看各个gpu设备的使用情况,可以在服务器上的jupyter主页点击new->terminal,在terminal中输入nvidia-smi即可查看每张卡的使用情况.如下图.



上图左边一栏显示了他们的设备id(0,1,2,3),风扇转速,温度,性能状态,能耗等信息,中间一栏显示他们的bus-id和显存使用量,右边一栏是GPU使用率等信息.注意到中间一栏的显存使用量,在训练模型前我们可以根据空余的显存来选择我们使用的gpu设备.

在本次实验中我们将代码中的torch.device('cuda:0')的0更换成所需的设备id即可选择在相应的gpu设备上运行程序.

## CNN(卷积神经网络)

### 一个简单的CNN

接下来,让我们建立一个简单的CNN分类器. 这个CNN的整体流程是

卷积(Conv2d) -> BN(batch normalization) -> 激励函数(ReLU) -> 池化(MaxPooling) ->  
 卷积(Conv2d) -> BN(batch normalization) -> 激励函数(ReLU) -> 池化(MaxPooling) ->  
 全连接层(Linear) -> 输出.

In [12]:

```

import torch
import torch.nn as nn
import torch.utils.data as Data
import torchvision

class MyCNN(nn.Module):

    def __init__(self, image_size, num_classes):
        super(MyCNN, self).__init__()
        # conv1: Conv2d -> BN -> ReLU -> MaxPool
        self.conv1 = nn.Sequential(
            nn.Conv2d(in_channels=3, out_channels=16, kernel_size=3, stride=1, padding=1),
            nn.BatchNorm2d(16),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size=2, stride=2),
        )
        # conv2: Conv2d -> BN -> ReLU -> MaxPool
        self.conv2 = nn.Sequential(
            nn.Conv2d(in_channels=16, out_channels=32, kernel_size=3, stride=1, padding=1),
            nn.BatchNorm2d(32),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size=2, stride=2),
        )
        # fully connected layer
        self.fc = nn.Linear(32 * (image_size // 4) * (image_size // 4), num_classes)

    def forward(self, x):
        """
        input: N * 3 * image_size * image_size
        output: N * num_classes
        """
        x = self.conv1(x)
        x = self.conv2(x)
        # view(x.size(0), -1): change tensor size from (N, H, W) to (N, H*W)
        x = x.view(x.size(0), -1)
        output = self.fc(x)
        return output

```

这样,一个简单的CNN模型就写好了.与前面的课堂内容相似,我们需要对完成网络进行训练与评估的代码.

In [13]:

```
def train(model, train_loader, loss_func, optimizer, device):  
    """  
    train model using loss_fn and optimizer in an epoch.  
    model: CNN networks  
    train_loader: a Dataloader object with training data  
    loss_func: loss function  
    device: train on cpu or gpu device  
    """  
  
    total_loss = 0  
    # train the model using minibatch  
    for i, (images, targets) in enumerate(train_loader):  
        images = images.to(device)  
        targets = targets.to(device)  
  
        # forward  
        outputs = model(images)  
        loss = loss_func(outputs, targets)  
  
        # backward and optimize  
        optimizer.zero_grad()  
        loss.backward()  
        optimizer.step()  
  
        total_loss += loss.item()  
  
        # every 100 iteration, print loss  
        if (i + 1) % 100 == 0:  
            print ("Step [{}/{}] Train Loss: {:.4f}"  
                  .format(i+1, len(train_loader), loss.item()))  
    return total_loss / len(train_loader)
```

In [14]:

```

def evaluate(model, val_loader, device):
    """
    model: CNN networks
    val_loader: a Dataloader object with validation data
    device: evaluate on cpu or gpu device
    return classification accuracy of the model on val dataset
    """

    # evaluate the model
    model.eval()
    # context-manager that disabled gradient computation
    with torch.no_grad():
        correct = 0
        total = 0

        for i, (images, targets) in enumerate(val_loader):
            # device: cpu or gpu
            images = images.to(device)
            targets = targets.to(device)

            outputs = model(images)

            # return the maximum value of each row of the input tensor in the
            # given dimension dim, the second return value is the index location
            # of each maximum value found(argmax)
            _, predicted = torch.max(outputs.data, dim=1)

            correct += (predicted == targets).sum().item()

            total += targets.size(0)

        accuracy = correct / total
        print('Accuracy on Test Set: {:.4f} %'.format(100 * accuracy))
        return accuracy

```

In [15]:

```

def save_model(model, save_path):
    # save model
    torch.save(model.state_dict(), save_path)

```

In [16]:

```
import matplotlib.pyplot as plt
def show_curve(ys, title):
    """
    plot curve for Loss and Accuracy
    Args:
        ys: loss or acc list
        title: loss or accuracy
    """
    x = np.array(range(len(ys)))
    y = np.array(ys)
    plt.plot(x, y, c='b')
    plt.axis()
    plt.title('{} curve'.format(title))
    plt.xlabel('epoch')
    plt.ylabel('{}'.format(title))
    plt.show()
```

## 准备数据与训练模型

接下来，我们使用CIFAR10数据集来对我们的CNN模型进行训练。

CIFAR-10:该数据集共有60000张彩色图像,这些图像是32\*32,分为10个类,每类6000张图.这里面有50000张用于训练,构成了5个训练批,每一批10000张图;另外10000用于测试,单独构成一批.在本次实验中,使用CIFAR-10数据集来训练我们的模型.我们可以用torchvision.datasets.CIFAR10来直接使用CIFAR10数据集.

 cifar10

In [17]:

```
import torch
import torch.nn as nn
import torchvision
import torchvision.transforms as transforms

# mean and std of cifar10 in 3 channels
cifar10_mean = (0.49, 0.48, 0.45)
cifar10_std = (0.25, 0.24, 0.26)

# define transform operations of train dataset
train_transform = transforms.Compose([
    # data augmentation
    transforms.Pad(4),
    transforms.RandomHorizontalFlip(),
    transforms.RandomCrop(32),

    transforms.ToTensor(),
    transforms.Normalize(cifar10_mean, cifar10_std)])

test_transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize(cifar10_mean, cifar10_std)])

# torchvision.datasets provide CIFAR-10 dataset for classification
train_dataset = torchvision.datasets.CIFAR10(root='./data/',
                                              train=True,
                                              transform=train_transform,
                                              download=True)

test_dataset = torchvision.datasets.CIFAR10(root='./data/',
                                              train=False,
                                              transform=test_transform)

# Data loader: provides single- or multi-process iterators over the dataset.
train_loader = torch.utils.data.DataLoader(dataset=train_dataset,
                                           batch_size=100,
                                           shuffle=True)

test_loader = torch.utils.data.DataLoader(dataset=test_dataset,
                                           batch_size=100,
                                           shuffle=False)
```

Files already downloaded and verified

训练过程中使用交叉熵(cross-entropy)损失函数与Adam优化器来训练我们的分类器网络. 阅读下面的代码并在To-Do处,根据之前所学的知识,补充前向传播和反向传播的代码来实现分类网络的训练.

In [18]:

```
def fit(model, num_epochs, optimizer, device):  
    """  
    train and evaluate an classifier num_epochs times.  
    We use optimizer and cross entropy loss to train the model.  
    Args:  
        model: CNN network  
        num_epochs: the number of training epochs  
        optimizer: optimize the loss function  
    """  
  
    # loss and optimizer  
    loss_func = nn.CrossEntropyLoss()  
  
    model.to(device)  
    loss_func.to(device)  
  
    # log train loss and test accuracy  
    losses = []  
    accs = []  
  
    for epoch in range(num_epochs):  
  
        print('Epoch {}/{}:'.format(epoch + 1, num_epochs))  
        # train step  
        loss = train(model, train_loader, loss_func, optimizer, device)  
        losses.append(loss)  
  
        # evaluate step  
        accuracy = evaluate(model, test_loader, device)  
        accs.append(accuracy)  
  
    # show curve  
    show_curve(losses, "train loss")  
    show_curve(accs, "test accuracy")
```

In [19]:

```
# hyper parameters  
num_epochs = 10  
lr = 0.01  
image_size = 32  
num_classes = 10
```



In [20]:

```
# declare and define an objet of MyCNN
mycnn = MyCNN(image_size, num_classes)
print(mycnn)
```

```
MyCNN(
  (conv1): Sequential(
    (0): Conv2d(3, 16, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1))
    (1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True, track
_running_stats=True)
    (2): ReLU()
    (3): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, c
eil_mode=False)
  )
  (conv2): Sequential(
    (0): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1))
    (1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track
_running_stats=True)
    (2): ReLU()
    (3): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, c
eil_mode=False)
  )
  (fc): Linear(in_features=2048, out_features=10, bias=True)
)
```

In [21]:

```
# Device configuration, cpu, cuda:0/1/2/3 available
device = torch.device('cuda:0')

optimizer = torch.optim.Adam(mycnn.parameters(), lr=lr)

# start training on cifar10 dataset
fit(mycnn, num_epochs, optimizer, device)
```

Epoch 1/10:

Step [100/500] Train Loss: 1.8075  
Step [200/500] Train Loss: 1.6811  
Step [300/500] Train Loss: 1.6177  
Step [400/500] Train Loss: 1.3389  
Step [500/500] Train Loss: 1.2736  
Accuracy on Test Set: 53.9500 %

Epoch 2/10:

Step [100/500] Train Loss: 1.5978  
Step [200/500] Train Loss: 1.2951  
Step [300/500] Train Loss: 1.3162  
Step [400/500] Train Loss: 1.2874  
Step [500/500] Train Loss: 1.1236  
Accuracy on Test Set: 61.5300 %

Epoch 3/10:

Step [100/500] Train Loss: 1.3468  
Step [200/500] Train Loss: 1.3069  
Step [300/500] Train Loss: 1.1912  
Step [400/500] Train Loss: 1.2451  
Step [500/500] Train Loss: 1.3067  
Accuracy on Test Set: 60.2800 %

Epoch 4/10:

Step [100/500] Train Loss: 1.3471  
Step [200/500] Train Loss: 1.2564  
Step [300/500] Train Loss: 1.1971  
Step [400/500] Train Loss: 1.1134  
Step [500/500] Train Loss: 1.3163  
Accuracy on Test Set: 62.7700 %

Epoch 5/10:

Step [100/500] Train Loss: 1.2081  
Step [200/500] Train Loss: 1.0366  
Step [300/500] Train Loss: 1.0514  
Step [400/500] Train Loss: 1.1292  
Step [500/500] Train Loss: 1.0381  
Accuracy on Test Set: 64.4700 %

Epoch 6/10:

Step [100/500] Train Loss: 0.9613  
Step [200/500] Train Loss: 0.9588  
Step [300/500] Train Loss: 1.1643  
Step [400/500] Train Loss: 0.9842  
Step [500/500] Train Loss: 1.0876  
Accuracy on Test Set: 64.2500 %

Epoch 7/10:

Step [100/500] Train Loss: 1.1227  
Step [200/500] Train Loss: 1.1365  
Step [300/500] Train Loss: 1.2146  
Step [400/500] Train Loss: 1.0229  
Step [500/500] Train Loss: 1.3981  
Accuracy on Test Set: 65.6000 %



Epoch 8/10:

Step [100/500] Train Loss: 1.1427  
Step [200/500] Train Loss: 0.9221  
Step [300/500] Train Loss: 1.1509  
Step [400/500] Train Loss: 0.9516  
Step [500/500] Train Loss: 1.1159  
Accuracy on Test Set: 65.5400 %

Epoch 9/10:

Step [100/500] Train Loss: 1.0614  
Step [200/500] Train Loss: 1.0258  
Step [300/500] Train Loss: 0.9749  
Step [400/500] Train Loss: 0.9400

```
Step [500/500] Train Loss: 1.2101
Accuracy on Test Set: 66.7200 %
Epoch 10/10:
Step [100/500] Train Loss: 1.2158
Step [200/500] Train Loss: 1.1549
Step [300/500] Train Loss: 0.9802
Step [400/500] Train Loss: 0.9733
Step [500/500] Train Loss: 1.0673
Accuracy on Test Set: 66.6800 %
```



## ResNet

接下来,让我们完成更复杂的CNN的实现.

ResNet又叫做残差网络.在ResNet网络结构中会用到两种残差模块,一种是以两个3\*3的卷积网络串接在一起作为一个残差模块,另外一种1\*1、3\*3、1\*1的3个卷积网络串接在一起作为一个残差模块.他们如下图所示.



我们以左边的模块为例实现一个ResidualBlock.注意到由于我们在两次卷积中可能会使输入的tensor的size与输出的tensor的size不相等,为了使它们能够相加,所以输出的tensor与输入的tensor size不同时,我们使用downsample(由外部传入)来使保持size相同

现在,试在**To-Do**补充代码完成下面的forward函数来完成ResidualBlock的实现,并运行它.

In [22]:

```
# 3x3 convolution
def conv3x3(in_channels, out_channels, stride=1):
    return nn.Conv2d(in_channels, out_channels, kernel_size=3,
                      stride=stride, padding=1, bias=False)

# Residual block
class ResidualBlock(nn.Module):
    def __init__(self, in_channels, out_channels, stride=1, downsample=None):
        super(ResidualBlock, self).__init__()
        self.conv1 = conv3x3(in_channels, out_channels, stride)
        self.bn1 = nn.BatchNorm2d(out_channels)
        self.relu = nn.ReLU(inplace=True)
        self.conv2 = conv3x3(out_channels, out_channels)
        self.bn2 = nn.BatchNorm2d(out_channels)
        self.downsample = downsample

    def forward(self, x):
        """
        Defines the computation performed at every call.
        x: N * C * H * W
        """
        residual = x
        # if the size of input x changes, using downsample to change the size of
        residual
        if self.downsample:
            residual = self.downsample(x)
        out = self.conv1(x)
        out = self.bn1(out)

        """
        To-Do: add code here
        """
        out = self.relu(out)
        out = self.conv2(out)
        out = self.bn2(out)

        out += residual
        out = self.relu(out)
        return out
```

下面是一份针对cifar10数据集的ResNet的实现.它先通过一个conv3x3,然后经过3个包含多个残差模块的layer(一个layer可能包括多个ResidualBlock, 由传入的layers列表中的数字决定), 然后经过一个全局平均池化层, 最后通过一个线性层.

In [23]:

```

class ResNet(nn.Module):
    def __init__(self, block, layers, num_classes=10):
        """
        block: ResidualBlock or other block
        layers: a list with 3 positive num.
        """
        super(ResNet, self).__init__()
        self.in_channels = 16
        self.conv = conv3x3(3, 16)
        self.bn = nn.BatchNorm2d(16)
        self.relu = nn.ReLU(inplace=True)
        # layer1: image size 32
        self.layer1 = self.make_layer(block, 16, num_blocks=layers[0])
        # layer2: image size 32 -> 16
        self.layer2 = self.make_layer(block, 32, num_blocks=layers[1], stride=2)
        # layer3: image size 16 -> 8
        self.layer3 = self.make_layer(block, 64, num_blocks=layers[2], stride=2)
        # global avg pool: image size 8 -> 1
        self.avg_pool = nn.AvgPool2d(8)

        self.fc = nn.Linear(64, num_classes)

    def make_layer(self, block, out_channels, num_blocks, stride=1):
        """
        make a layer with num_blocks blocks.
        """
        downsample = None
        if (stride != 1) or (self.in_channels != out_channels):
            # use Conv2d with stride to downsample
            downsample = nn.Sequential(
                conv3x3(self.in_channels, out_channels, stride=stride),
                nn.BatchNorm2d(out_channels))

        # first block with downsample
        layers = []
        layers.append(block(self.in_channels, out_channels, stride, downsample))

        self.in_channels = out_channels
        # add num_blocks - 1 blocks
        for i in range(1, num_blocks):
            layers.append(block(out_channels, out_channels))

        # return a layer containing layers
        return nn.Sequential(*layers)

    def forward(self, x):
        out = self.conv(x)
        out = self.bn(out)
        out = self.relu(out)
        out = self.layer1(out)
        out = self.layer2(out)
        out = self.layer3(out)
        out = self.avg_pool(out)
        # view: here change output size from 4 dimensions to 2 dimensions
        out = out.view(out.size(0), -1)
        out = self.fc(out)
        return out

```





In [24]:

```
resnet = ResNet(ResidualBlock, [2, 2, 2])  
print(resnet)
```

```

ResNet(
  (conv): Conv2d(3, 16, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1), bias=False)
  (bn): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True, track_
running_stats=True)
  (relu): ReLU(inplace)
  (layer1): Sequential(
    (0): ResidualBlock(
      (conv1): Conv2d(16, 16, kernel_size=(3, 3), stride=(1, 1), pad
ding=(1, 1), bias=False)
      (bn1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True, t
rack_running_stats=True)
      (relu): ReLU(inplace)
      (conv2): Conv2d(16, 16, kernel_size=(3, 3), stride=(1, 1), pad
ding=(1, 1), bias=False)
      (bn2): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True, t
rack_running_stats=True)
    )
    (1): ResidualBlock(
      (conv1): Conv2d(16, 16, kernel_size=(3, 3), stride=(1, 1), pad
ding=(1, 1), bias=False)
      (bn1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True, t
rack_running_stats=True)
      (relu): ReLU(inplace)
      (conv2): Conv2d(16, 16, kernel_size=(3, 3), stride=(1, 1), pad
ding=(1, 1), bias=False)
      (bn2): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True, t
rack_running_stats=True)
    )
  )
  (layer2): Sequential(
    (0): ResidualBlock(
      (conv1): Conv2d(16, 32, kernel_size=(3, 3), stride=(2, 2), pad
ding=(1, 1), bias=False)
      (bn1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, t
rack_running_stats=True)
      (relu): ReLU(inplace)
      (conv2): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1), pad
ding=(1, 1), bias=False)
      (bn2): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, t
rack_running_stats=True)
      (downsample): Sequential(
        (0): Conv2d(16, 32, kernel_size=(3, 3), stride=(2, 2), paddi
ng=(1, 1), bias=False)
        (1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, t
rack_running_stats=True)
      )
    )
    (1): ResidualBlock(
      (conv1): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1), pad
ding=(1, 1), bias=False)
      (bn1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, t
rack_running_stats=True)
      (relu): ReLU(inplace)
      (conv2): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1), pad
ding=(1, 1), bias=False)
      (bn2): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, t
rack_running_stats=True)
    )
  )
  (layer3): Sequential(

```

```

(0): ResidualBlock(
  (conv1): Conv2d(32, 64, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
  (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace)
  (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (downsample): Sequential(
    (0): Conv2d(32, 64, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
    (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  )
)
(1): ResidualBlock(
  (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace)
  (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
)
)
(avg_pool): AvgPool2d(kernel_size=8, stride=8, padding=0)
(fc): Linear(in_features=64, out_features=10, bias=True)
)

```

使用fit函数训练实现的ResNet,观察结果变化.



In [26]:

```
# Hyper-parameters
num_epochs = 10
lr = 0.001
# Device configuration
device = torch.device('cuda:0')
# optimizer
optimizer = torch.optim.Adam(resnet.parameters(), lr=lr)

fit(resnet, num_epochs, optimizer, device)
```

Epoch 1/10:  
Step [100/500] Train Loss: 1.0425  
Step [200/500] Train Loss: 1.2821  
Step [300/500] Train Loss: 1.0189  
Step [400/500] Train Loss: 1.0343  
Step [500/500] Train Loss: 1.0760  
Accuracy on Test Set: 63.9400 %  
Epoch 2/10:  
Step [100/500] Train Loss: 0.9691  
Step [200/500] Train Loss: 0.9280  
Step [300/500] Train Loss: 1.1253  
Step [400/500] Train Loss: 1.0832  
Step [500/500] Train Loss: 0.7534  
Accuracy on Test Set: 63.9400 %  
Epoch 3/10:  
Step [100/500] Train Loss: 0.9576  
Step [200/500] Train Loss: 0.8765  
Step [300/500] Train Loss: 0.7416  
Step [400/500] Train Loss: 0.8020  
Step [500/500] Train Loss: 0.7128  
Accuracy on Test Set: 68.0000 %  
Epoch 4/10:  
Step [100/500] Train Loss: 1.0099  
Step [200/500] Train Loss: 0.9608  
Step [300/500] Train Loss: 0.8774  
Step [400/500] Train Loss: 0.7870  
Step [500/500] Train Loss: 0.7058  
Accuracy on Test Set: 68.5800 %  
Epoch 5/10:  
Step [100/500] Train Loss: 0.8077  
Step [200/500] Train Loss: 0.5876  
Step [300/500] Train Loss: 0.8926  
Step [400/500] Train Loss: 0.8441  
Step [500/500] Train Loss: 0.9973  
Accuracy on Test Set: 72.6900 %  
Epoch 6/10:  
Step [100/500] Train Loss: 0.8229  
Step [200/500] Train Loss: 0.7058  
Step [300/500] Train Loss: 0.7750  
Step [400/500] Train Loss: 0.7295  
Step [500/500] Train Loss: 0.8246  
Accuracy on Test Set: 72.6600 %  
Epoch 7/10:  
Step [100/500] Train Loss: 0.7068  
Step [200/500] Train Loss: 0.6928  
Step [300/500] Train Loss: 0.8502  
Step [400/500] Train Loss: 0.7325  
Step [500/500] Train Loss: 0.6583  
Accuracy on Test Set: 75.1100 %  
Epoch 8/10:  
Step [100/500] Train Loss: 0.6834  
Step [200/500] Train Loss: 0.8615  
Step [300/500] Train Loss: 0.7363  
Step [400/500] Train Loss: 0.8829  
Step [500/500] Train Loss: 0.7208  
Accuracy on Test Set: 74.1100 %  
Epoch 9/10:  
Step [100/500] Train Loss: 0.6611  
Step [200/500] Train Loss: 0.5346  
Step [300/500] Train Loss: 0.4550  
Step [400/500] Train Loss: 0.7190

```
Step [500/500] Train Loss: 0.5672
Accuracy on Test Set: 76.9400 %
Epoch 10/10:
Step [100/500] Train Loss: 0.5207
Step [200/500] Train Loss: 0.6895
Step [300/500] Train Loss: 0.5880
Step [400/500] Train Loss: 0.6893
Step [500/500] Train Loss: 0.7157
Accuracy on Test Set: 77.9500 %
```



## 作业

尝试改变学习率 $lr$ ,使用SGD或Adam优化器,训练10个epoch,提高ResNet在测试集上的accuracy.

In [30]:



```
# Hyper-parameters
num_epochs = 10
lr = 0.0015
# Device configuration
device = torch.device('cuda:0')
# optimizer
optimizer = torch.optim.Adam(resnet.parameters(), lr=lr)

fit(resnet, num_epochs, optimizer, device)
```

Epoch 1/10:  
Step [100/500] Train Loss: 0.7118  
Step [200/500] Train Loss: 0.4573  
Step [300/500] Train Loss: 0.4669  
Step [400/500] Train Loss: 0.2568  
Step [500/500] Train Loss: 0.4969  
Accuracy on Test Set: 80.3800 %  
Epoch 2/10:  
Step [100/500] Train Loss: 0.4439  
Step [200/500] Train Loss: 0.4941  
Step [300/500] Train Loss: 0.5434  
Step [400/500] Train Loss: 0.4898  
Step [500/500] Train Loss: 0.4460  
Accuracy on Test Set: 82.1700 %  
Epoch 3/10:  
Step [100/500] Train Loss: 0.4875  
Step [200/500] Train Loss: 0.3971  
Step [300/500] Train Loss: 0.5229  
Step [400/500] Train Loss: 0.6836  
Step [500/500] Train Loss: 0.4133  
Accuracy on Test Set: 78.1500 %  
Epoch 4/10:  
Step [100/500] Train Loss: 0.3835  
Step [200/500] Train Loss: 0.5045  
Step [300/500] Train Loss: 0.4055  
Step [400/500] Train Loss: 0.3561  
Step [500/500] Train Loss: 0.4818  
Accuracy on Test Set: 83.5100 %  
Epoch 5/10:  
Step [100/500] Train Loss: 0.3647  
Step [200/500] Train Loss: 0.5745  
Step [300/500] Train Loss: 0.2970  
Step [400/500] Train Loss: 0.4631  
Step [500/500] Train Loss: 0.3952  
Accuracy on Test Set: 82.9100 %  
Epoch 6/10:  
Step [100/500] Train Loss: 0.4992  
Step [200/500] Train Loss: 0.4990  
Step [300/500] Train Loss: 0.4383  
Step [400/500] Train Loss: 0.5731  
Step [500/500] Train Loss: 0.3213  
Accuracy on Test Set: 83.0500 %  
Epoch 7/10:  
Step [100/500] Train Loss: 0.3208  
Step [200/500] Train Loss: 0.3100  
Step [300/500] Train Loss: 0.4275  
Step [400/500] Train Loss: 0.4537  
Step [500/500] Train Loss: 0.4117  
Accuracy on Test Set: 83.2300 %  
Epoch 8/10:  
Step [100/500] Train Loss: 0.4122  
Step [200/500] Train Loss: 0.4852  
Step [300/500] Train Loss: 0.4390  
Step [400/500] Train Loss: 0.3829  
Step [500/500] Train Loss: 0.3836  
Accuracy on Test Set: 83.1100 %  
Epoch 9/10:  
Step [100/500] Train Loss: 0.3871  
Step [200/500] Train Loss: 0.3587  
Step [300/500] Train Loss: 0.2804  
Step [400/500] Train Loss: 0.2926




```
Step [500/500] Train Loss: 0.4059
Accuracy on Test Set: 83.7800 %
Epoch 10/10:
Step [100/500] Train Loss: 0.3101
Step [200/500] Train Loss: 0.4478
Step [300/500] Train Loss: 0.3073
Step [400/500] Train Loss: 0.3947
Step [500/500] Train Loss: 0.3530
Accuracy on Test Set: 84.1200 %
```



## 作业

下图表示将SE模块嵌入到ResNet的残差模块.

 SE-Resnet module

其中,global pooling表示全局池化层(将输入的size池化为 $1 \times 1$ ), 将 $c \times h \times w$ 的输入变为 $c \times 1 \times 1$ 的输出.FC表示全连接层(线性层),两层FC之间使用ReLU作为激活函数.通过两层FC后使用sigmoid激活函数激活.最后将得到的 $c$ 个值与原输入 $c \times h \times w$ 按channel相乘,得到 $c \times h \times w$ 的输出.

补充下方的代码完成SE-Resnet block的实现.

In [31]:

```

class SELayer(nn.Module):
    def __init__(self, channel, reduction=16):
        super(SELayer, self).__init__()
        # The output of AdaptiveAvgPool2d is of size H x W, for any input size.
        self.avg_pool = nn.AdaptiveAvgPool2d((1, 1))
        self.relu = nn.ReLU(inplace=True)
        self.fc1 = nn.Linear(channel, channel//reduction)
        self.fc2 = nn.Linear(channel//reduction, channel)
        self.sigmoid = nn.Sigmoid()

    def forward(self, x):
        out = self.avg_pool(x)
        out = out.view(out.size(0), -1)
        out = self.fc1(out)
        out = self.relu(out)
        out = self.fc2(out)
        out = self.sigmoid(out)
        out = out.view(out.shape[0], -1, 1, 1)
        return x*out

```

In [32]:

```

class SEResidualBlock(nn.Module):
    def __init__(self, in_channels, out_channels, stride=1, downsample=None, reduction=16):
        super(SEResidualBlock, self).__init__()
        """
        To-Do: add code here
        """

        self.conv1 = conv3x3(in_channels, out_channels, stride)
        self.bn1 = nn.BatchNorm2d(out_channels)
        self.relu = nn.ReLU(inplace=True)
        self.conv2 = conv3x3(out_channels, out_channels)
        self.bn2 = nn.BatchNorm2d(out_channels)
        self.se = SELayer(out_channels, reduction)
        self.downsample = downsample

    def forward(self, x):

        residual = x
        """
        To-Do: add code here
        """

        if self.downsample:
            residual = self.downsample(x)
        out = self.conv1(x)
        out = self.bn1(out)
        out = self.relu(out)
        out = self.conv2(out)
        out = self.bn2(out)
        out = self.se(out)
        out = out + residual
        out = self.relu(out)
        return out

```

In [33]:

```
se_resnet = ResNet(SEResidualBlock, [2, 2, 2])  
print(se_resnet)
```

```

ResNet(
  (conv): Conv2d(3, 16, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1), bias=False)
  (bn): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True, track_
running_stats=True)
  (relu): ReLU(inplace)
  (layer1): Sequential(
    (0): SEResidualBlock(
      (conv1): Conv2d(16, 16, kernel_size=(3, 3), stride=(1, 1), pad
ding=(1, 1), bias=False)
      (bn1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True, t
rack_running_stats=True)
      (relu): ReLU(inplace)
      (conv2): Conv2d(16, 16, kernel_size=(3, 3), stride=(1, 1), pad
ding=(1, 1), bias=False)
      (bn2): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True, t
rack_running_stats=True)
      (se): SELayer(
        (avg_pool): AdaptiveAvgPool2d(output_size=(1, 1))
        (relu): ReLU(inplace)
        (fc1): Linear(in_features=16, out_features=1, bias=True)
        (fc2): Linear(in_features=1, out_features=16, bias=True)
        (sigmoid): Sigmoid()
      )
    )
    (1): SEResidualBlock(
      (conv1): Conv2d(16, 16, kernel_size=(3, 3), stride=(1, 1), pad
ding=(1, 1), bias=False)
      (bn1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True, t
rack_running_stats=True)
      (relu): ReLU(inplace)
      (conv2): Conv2d(16, 16, kernel_size=(3, 3), stride=(1, 1), pad
ding=(1, 1), bias=False)
      (bn2): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True, t
rack_running_stats=True)
      (se): SELayer(
        (avg_pool): AdaptiveAvgPool2d(output_size=(1, 1))
        (relu): ReLU(inplace)
        (fc1): Linear(in_features=16, out_features=1, bias=True)
        (fc2): Linear(in_features=1, out_features=16, bias=True)
        (sigmoid): Sigmoid()
      )
    )
  )
  (layer2): Sequential(
    (0): SEResidualBlock(
      (conv1): Conv2d(16, 32, kernel_size=(3, 3), stride=(2, 2), pad
ding=(1, 1), bias=False)
      (bn1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, t
rack_running_stats=True)
      (relu): ReLU(inplace)
      (conv2): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1), pad
ding=(1, 1), bias=False)
      (bn2): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, t
rack_running_stats=True)
      (se): SELayer(
        (avg_pool): AdaptiveAvgPool2d(output_size=(1, 1))
        (relu): ReLU(inplace)
        (fc1): Linear(in_features=32, out_features=2, bias=True)
        (fc2): Linear(in_features=2, out_features=32, bias=True)
        (sigmoid): Sigmoid()
      )
    )
  )
)

```

```

    )
    (downsample): Sequential(
      (0): Conv2d(16, 32, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
      (1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
  (1): SEResidualBlock(
    (conv1): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace)
    (conv2): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (se): SELayer(
      (avg_pool): AdaptiveAvgPool2d(output_size=(1, 1))
      (relu): ReLU(inplace)
      (fc1): Linear(in_features=32, out_features=2, bias=True)
      (fc2): Linear(in_features=2, out_features=32, bias=True)
      (sigmoid): Sigmoid()
    )
  )
)
(layer3): Sequential(
  (0): SEResidualBlock(
    (conv1): Conv2d(32, 64, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace)
    (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (se): SELayer(
      (avg_pool): AdaptiveAvgPool2d(output_size=(1, 1))
      (relu): ReLU(inplace)
      (fc1): Linear(in_features=64, out_features=4, bias=True)
      (fc2): Linear(in_features=4, out_features=64, bias=True)
      (sigmoid): Sigmoid()
    )
  )
  (downsample): Sequential(
    (0): Conv2d(32, 64, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
    (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  )
  (1): SEResidualBlock(
    (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace)
    (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  )
)

```

```
rack_running_stats=True)
    (se): SELayer(
      (avg_pool): AdaptiveAvgPool2d(output_size=(1, 1))
      (relu): ReLU(inplace)
      (fc1): Linear(in_features=64, out_features=4, bias=True)
      (fc2): Linear(in_features=4, out_features=64, bias=True)
      (sigmoid): Sigmoid()
    )
  )
)
(avg_pool): AvgPool2d(kernel_size=8, stride=8, padding=0)
(fc): Linear(in_features=64, out_features=10, bias=True)
)
```

In [34]:



```
# Hyper-parameters
num_epochs = 10
lr = 0.001
# Device configuration
device = torch.device('cuda:0')
# optimizer
optimizer = torch.optim.Adam(se_resnet.parameters(), lr=lr)

fit(se_resnet, num_epochs, optimizer, device)
```

Epoch 1/10:  
Step [100/500] Train Loss: 1.6276  
Step [200/500] Train Loss: 1.4714  
Step [300/500] Train Loss: 1.4851  
Step [400/500] Train Loss: 1.2222  
Step [500/500] Train Loss: 1.2060  
Accuracy on Test Set: 48.9400 %  
Epoch 2/10:  
Step [100/500] Train Loss: 2.2510  
Step [200/500] Train Loss: 2.0723  
Step [300/500] Train Loss: 1.8598  
Step [400/500] Train Loss: 2.0755  
Step [500/500] Train Loss: 1.7243  
Accuracy on Test Set: 33.7100 %  
Epoch 3/10:  
Step [100/500] Train Loss: 1.7078  
Step [200/500] Train Loss: 1.5886  
Step [300/500] Train Loss: 1.5629  
Step [400/500] Train Loss: 1.5738  
Step [500/500] Train Loss: 1.4202  
Accuracy on Test Set: 48.1800 %  
Epoch 4/10:  
Step [100/500] Train Loss: 1.5383  
Step [200/500] Train Loss: 1.4838  
Step [300/500] Train Loss: 1.3516  
Step [400/500] Train Loss: 1.4415  
Step [500/500] Train Loss: 1.1955  
Accuracy on Test Set: 54.4100 %  
Epoch 5/10:  
Step [100/500] Train Loss: 1.2495  
Step [200/500] Train Loss: 1.2082  
Step [300/500] Train Loss: 1.1445  
Step [400/500] Train Loss: 1.0991  
Step [500/500] Train Loss: 1.1674  
Accuracy on Test Set: 56.0800 %  
Epoch 6/10:  
Step [100/500] Train Loss: 1.0126  
Step [200/500] Train Loss: 1.1029  
Step [300/500] Train Loss: 0.8674  
Step [400/500] Train Loss: 0.9355  
Step [500/500] Train Loss: 1.1729  
Accuracy on Test Set: 61.1100 %  
Epoch 7/10:  
Step [100/500] Train Loss: 1.1173  
Step [200/500] Train Loss: 1.2414  
Step [300/500] Train Loss: 1.1263  
Step [400/500] Train Loss: 1.0653  
Step [500/500] Train Loss: 0.9470  
Accuracy on Test Set: 61.7000 %  
Epoch 8/10:  
Step [100/500] Train Loss: 1.0067  
Step [200/500] Train Loss: 0.9689  
Step [300/500] Train Loss: 0.9487  
Step [400/500] Train Loss: 1.1266  
Step [500/500] Train Loss: 1.1523  
Accuracy on Test Set: 66.2600 %  
Epoch 9/10:  
Step [100/500] Train Loss: 0.7574  
Step [200/500] Train Loss: 0.7837  
Step [300/500] Train Loss: 0.9518  
Step [400/500] Train Loss: 0.9028



Step [500/500] Train Loss: 0.8175  
Accuracy on Test Set: 66.4400 %  
Epoch 10/10:  
Step [100/500] Train Loss: 0.7346  
Step [200/500] Train Loss: 0.7445  
Step [300/500] Train Loss: 0.8594  
Step [400/500] Train Loss: 0.9784  
Step [500/500] Train Loss: 0.8334  
Accuracy on Test Set: 67.4600 %



## Vgg

接下来让我们阅读vgg网络的实现代码。VGGNet全部使用3\*3的卷积核和2\*2的池化核，通过不断加深网络结构来提升性能。Vgg表明了卷积神经网络的深度增加和小卷积核的使用对网络的最终分类识别效果有很大的作用。



下面是一份用于训练cifar10的简化版的vgg代码。

有时间的同学可以阅读并训练它。

In [35]:

```
import math

class VGG(nn.Module):
    def __init__(self, cfg):
        super(VGG, self).__init__()
        self.features = self._make_layers(cfg)
        # linear layer
        self.classifier = nn.Linear(512, 10)

    def forward(self, x):
        out = self.features(x)
        out = out.view(out.size(0), -1)
        out = self.classifier(out)
        return out

    def _make_layers(self, cfg):
        """
        cfg: a list define layers this layer contains
        'M': MaxPool, number: Conv2d(out_channels=number) -> BN -> ReLU
        """
        layers = []
        in_channels = 3
        for x in cfg:
            if x == 'M':
                layers += [nn.MaxPool2d(kernel_size=2, stride=2)]
            else:
                layers += [nn.Conv2d(in_channels, x, kernel_size=3, padding=1),
                           nn.BatchNorm2d(x),
                           nn.ReLU(inplace=True)]
                in_channels = x
        layers += [nn.AvgPool2d(kernel_size=1, stride=1)]
        return nn.Sequential(*layers)
```

In [36]:

```
cfg = {  
    'VGG11': [64, 'M', 128, 'M', 256, 256, 'M', 512, 512, 'M', 512, 512, 'M'],  
    'VGG13': [64, 64, 'M', 128, 128, 'M', 256, 256, 'M', 512, 512, 'M', 512, 512  
, 'M'],  
    'VGG16': [64, 64, 'M', 128, 128, 'M', 256, 256, 256, 'M', 512, 512, 512, 'M'  
, 512, 512, 512, 'M'],  
    'VGG19': [64, 64, 'M', 128, 128, 'M', 256, 256, 256, 256, 'M', 512, 512, 512  
, 512, 'M', 512, 512, 512, 'M'],  
}  
vggnet = VGG(cfg['VGG11'])  
print(vggnet)
```

```

VGG(
  (features): Sequential(
    (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1))
    (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track
_running_stats=True)
    (2): ReLU(inplace)
    (3): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, c
eil_mode=False)
    (4): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1))
    (5): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, trac
k_running_stats=True)
    (6): ReLU(inplace)
    (7): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, c
eil_mode=False)
    (8): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding
=(1, 1))
    (9): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, trac
k_running_stats=True)
    (10): ReLU(inplace)
    (11): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), paddin
g=(1, 1))
    (12): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, tra
ck_running_stats=True)
    (13): ReLU(inplace)
    (14): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
    (15): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), paddin
g=(1, 1))
    (16): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, tra
ck_running_stats=True)
    (17): ReLU(inplace)
    (18): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), paddin
g=(1, 1))
    (19): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, tra
ck_running_stats=True)
    (20): ReLU(inplace)
    (21): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
    (22): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), paddin
g=(1, 1))
    (23): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, tra
ck_running_stats=True)
    (24): ReLU(inplace)
    (25): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), paddin
g=(1, 1))
    (26): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, tra
ck_running_stats=True)
    (27): ReLU(inplace)
    (28): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
    (29): AvgPool2d(kernel_size=1, stride=1, padding=0)
  )
  (classifier): Linear(in_features=512, out_features=10, bias=True)
)

```

In [37]:

```
# Hyper-parameters
num_epochs = 10
lr = 1e-3
# Device configuration
device = torch.device('cuda:0')

# optimizer
optimizer = torch.optim.Adam(vggnet.parameters(), lr=lr)

fit(vggnet, num_epochs, optimizer, device)
```

Epoch 1/10:  
Step [100/500] Train Loss: 1.6253  
Step [200/500] Train Loss: 1.4231  
Step [300/500] Train Loss: 1.3688  
Step [400/500] Train Loss: 1.3814  
Step [500/500] Train Loss: 0.9911  
Accuracy on Test Set: 57.4000 %  
Epoch 2/10:  
Step [100/500] Train Loss: 1.8048  
Step [200/500] Train Loss: 1.4972  
Step [300/500] Train Loss: 1.3364  
Step [400/500] Train Loss: 1.2925  
Step [500/500] Train Loss: 1.1823  
Accuracy on Test Set: 58.4400 %  
Epoch 3/10:  
Step [100/500] Train Loss: 1.1463  
Step [200/500] Train Loss: 0.9488  
Step [300/500] Train Loss: 1.1180  
Step [400/500] Train Loss: 0.9506  
Step [500/500] Train Loss: 0.8822  
Accuracy on Test Set: 69.1200 %  
Epoch 4/10:  
Step [100/500] Train Loss: 0.9562  
Step [200/500] Train Loss: 0.7132  
Step [300/500] Train Loss: 0.7834  
Step [400/500] Train Loss: 0.9923  
Step [500/500] Train Loss: 0.6245  
Accuracy on Test Set: 74.0900 %  
Epoch 5/10:  
Step [100/500] Train Loss: 0.6804  
Step [200/500] Train Loss: 0.7942  
Step [300/500] Train Loss: 0.6620  
Step [400/500] Train Loss: 0.5886  
Step [500/500] Train Loss: 0.6147  
Accuracy on Test Set: 78.1000 %  
Epoch 6/10:  
Step [100/500] Train Loss: 0.4513  
Step [200/500] Train Loss: 0.6562  
Step [300/500] Train Loss: 0.5617  
Step [400/500] Train Loss: 0.6486  
Step [500/500] Train Loss: 0.6400  
Accuracy on Test Set: 78.4500 %  
Epoch 7/10:  
Step [100/500] Train Loss: 0.6970  
Step [200/500] Train Loss: 0.5626  
Step [300/500] Train Loss: 0.4481  
Step [400/500] Train Loss: 0.5924  
Step [500/500] Train Loss: 0.5008  
Accuracy on Test Set: 80.9900 %  
Epoch 8/10:  
Step [100/500] Train Loss: 0.5288  
Step [200/500] Train Loss: 0.4491  
Step [300/500] Train Loss: 0.5524  
Step [400/500] Train Loss: 0.5024  
Step [500/500] Train Loss: 0.4200  
Accuracy on Test Set: 81.3000 %  
Epoch 9/10:  
Step [100/500] Train Loss: 0.5242  
Step [200/500] Train Loss: 0.4221  
Step [300/500] Train Loss: 0.4665  
Step [400/500] Train Loss: 0.6280

Step [500/500] Train Loss: 0.5573  
Accuracy on Test Set: 81.2000 %  
Epoch 10/10:  
Step [100/500] Train Loss: 0.3493  
Step [200/500] Train Loss: 0.5310  
Step [300/500] Train Loss: 0.6748  
Step [400/500] Train Loss: 0.4147  
Step [500/500] Train Loss: 0.4272  
Accuracy on Test Set: 83.5300 %

